

---

**University of Puerto Rico  
Campus of Mayagüez  
Mayagüez, Puerto Rico  
Department of Electrical and Computer Engineering**



**ARTIFICIAL INTELLIGENCE  
Group Report: Team H  
AI and Society: Ethical Impact Analysis**

**By**

**Claudia Sánchez Merino  
Manuel A. Infante Degro  
Carlos M. Chamocho Canals  
Emmanuel A. Morales Martínez**

**For: Prof. José F. Vega Riveros  
ICOM5015-S001D  
25 / May / 2026**

---

Table of contents:

1. 1. Introduction
2. Theoretical Background
  - a. Consequentialism
  - b. Deontological Ethics
  - c. Virtue Ethics
3. AI's Impact on Employment and Labor Markets
  - 3.1 Consequentialist Perspective
  - 3.2 Deontological Perspective
  - 3.3 Virtue Ethics Perspective
4. Privacy, Surveillance, and Civil Liberties
  - a. Consequentialist Perspective
  - b. Deontological Perspective
  - c. Virtue Ethics Perspective
5. Positive Impacts of AI in Healthcare and Science
  - 6.1 Positive Impacts in Healthcare
  - 6.2 Positive Impacts in Scientific Research
  - 6.3 Ethical Analysis
6. Group Deliberation Process
7. Conclusion
8. Future Work
9. References
10. Appendix A: Meeting Dates and Minutes
11. Appendix B: Individual Reflections

## **Introduction**

Artificial Intelligence is fundamentally transforming global society, reshaping all types of industries, from healthcare and education to surveillance and digital communication. Driven by rapid advancements in neural networks and deep learning, modern AI systems such as ChatGPT, Google Gemini, and DALL-E offer unprecedented capabilities. However, alongside their vast potential for innovation, these technologies introduce profound ethical challenges. As AI becomes deeply integrated into daily life, it is crucial to critically evaluate both its societal benefits and its inherent risks.

The purpose of this report is to conduct a comprehensive ethical analysis of AI's development and deployment by examining four critical areas of societal impact. First, it explores AI's disruptive effects on employment and labor markets. Second, it assesses the potential for malicious misuse through deepfakes, impersonation, and misinformation. Third, it examines the profound implications for privacy, surveillance, and civil liberties. Finally, it balances these concerns by highlighting AI's positive impacts, such as advancements in medical diagnostics, education, and scientific research.

To navigate these complex multi stakeholder issues, this analysis applies three major philosophical schools of ethics: consequentialism, deontological ethics, and virtue ethics. By utilizing these distinct ethical lenses, our group evaluates each of the four chosen domains to assess the moral implications of AI beyond its mere technological capability. Following individual research and analysis, our group convened to deliberate on these findings, exploring the overlaps, tensions, and conflicts between the different ethical perspectives to formulate the cohesive evaluation presented in this report

## **AI's Impact on Employment and Labor Markets**

The integration of Artificial Intelligence into the workforce is fundamentally altering global labor markets. While AI drives unprecedented productivity and economic growth, it also accelerates the automation of both manual and cognitive tasks. According to the World Economic Forum, macroeconomic trends and technology adoption will drive the creation of 69 million new jobs by 2027, but also the elimination of 83 million, resulting in a net decrease of 14 million jobs [1]. Furthermore, McKinsey Global Institute estimates that generative AI alone could enable the automation of up to 30% of hours currently worked across the US economy by 2030 [2]. This massive transition presents significant ethical dilemmas regarding worker displacement, economic inequality, and corporate responsibility.

From a consequentialist viewpoint, the ethicality of AI in the workplace is determined by its overall outcomes. A utilitarian would weigh the massive economic benefits such as increased corporate efficiency, lower costs of goods and services, and the creation of highly specialized new industries against the negative impacts of job losses and widened income inequality. If the wealth generated by AI is properly redistributed and the long-term standard of living increases for the majority, a consequentialist would likely view AI automation as morally right. However, if AI adoption disproportionately enriches a small minority of technology owners while thrusting millions into poverty, the action fails the utilitarian test of providing the greatest good for the greatest number.

For the optimistic case, an example can be seen from Goldman Sachs. Goldman Sachs estimates that AI adoption could raise productivity growth by approximately 1.5 percentage points per year over a ten-year period and lift global GDP by roughly 7%, or about \$7 trillion in additional output [7]. McKinsey projects that generative AI alone could enable the automation of up to 30% of hours currently worked across the U.S. economy by 2030, freeing labor for higher-value activity [2]. If these gains are realized and broadly distributed, a strict utilitarian calculus would favor accelerated AI deployment: the welfare of the many is increased, the cost of essential goods falls, and historical precedent suggests new industries eventually absorb displaced workers.

The pessimistic case, however, is supported by equally rigorous economic analysis and undermines the assumption that aggregate gains will translate into broad welfare improvement. Nobel laureate Daron Acemoglu's 2025 *Economic Policy* paper *The Simple*

*Macroeconomics of AI* models the labor-market effects of generative AI using a task-based framework and reaches notably sober conclusions: total factor productivity gains over the next ten years are estimated at less than 0.53%, far below industry projections, and Acemoglu demonstrates theoretically that even when AI improves the productivity of low-skill workers, the gains can flow disproportionately to capital owners and *increase* rather than decrease income inequality [8]. Acemoglu and Johnson have argued separately that on the current trajectory, the first-order impact of AI is likely to be increased inequality within industrial countries, with middle-income and lower-income countries also at significant risk [9]. Empirical analysis published in *Technology in Society* finds that AI capital accumulation is associated with worsening wealth inequality in a mutually reinforcing cycle, with income disparity amplifying wealth disparity and vice versa [10].

Deontology shifts the focus away from economic outcomes and toward the moral duties of employers and the inherent rights of workers. According to Kantian principles, humans must be treated as ends in themselves, never merely as means to an end. From this lens, mass layoffs executed solely to maximize shareholder profit via AI automation violate the employer's moral duty to their workforce. Furthermore, deontologists emphasize transparency and fairness. The use of "algorithmic management" where AI systems are used to hire, evaluate, or fire employees raises critical concerns about an individual's right to due process and an explanation, as workers may be penalized by opaque algorithms without human oversight.

Deontology imposes specific duties of transparency, fairness, and due process that current AI labor practices routinely violate. The most acute concerns center on what is called algorithmic management, where AI systems are used to hire, evaluate, discipline, and terminate workers without meaningful human judgment. Amazon's warehouse operations provide the clearest case. Internal documents disclosed by The Verge in 2019 revealed that Amazon's system "tracks the rates of each individual associate's productivity, and automatically generates any warnings or terminations regarding quality or productivity without input from supervisors" [11]. Workers have been terminated for productivity scores dipping below algorithmic thresholds with no opportunity to explain the context that a human manager would have considered routine [12]. By 2025, Amazon's algorithmic systems extended beyond warehouse floors: an October 2025 memo from senior vice president Beth Galetti announced layoffs explicitly designed to make the company "leaner and more agile"

by replacing middle-management functions with generative AI [13]. The pattern is now industry-wide. According to the Society for Human Resource Management, roughly one in four employers uses AI in HR practices, with 64% of those organizations deploying it in talent acquisition; 492 of the Fortune 500 use applicant tracking systems with AI components [14, 15].

Rather than looking at rules or outcomes, virtue ethics examines the character and intentions of the agents developing and deploying AI. A virtue ethicist asks: *Does the corporation's approach to AI automation reflect virtues like compassion, fairness, and responsibility, or vices like greed and apathy?* A virtuous organization would approach AI integration as a collaborative transition. Instead of treating employees as disposable assets, a morally responsible company would exhibit the virtue of care by investing in comprehensive upskilling and reskilling programs, ensuring that their workforce evolves alongside the technology. An example of an unjust application of AI in a corporation would be based on current examples and evidence based on the findings from The Bureau of Labor Statistics and Challenger, Gray & Christmas, where layoff announcements tracked were caused explicitly by AI. A figure of around 54,836 U.S. jobs in 2025 cited AI as its main cause for layoffs.. Independent econometric analyses using productivity and aggregate employment data estimate the true figure at roughly 200,000 to 300,000 U.S. jobs displaced or foregone due to AI in 2025 alone, four to six times the employer-reported number, with the gap concealed behind euphemisms like "restructuring" [3]. In the first six months of 2025, 77,999 tech-sector job losses were directly attributed to AI [4]. The World Economic Forum's 2025 Future of Jobs Report found 41% of employers worldwide intend to reduce their workforce within five years specifically because of AI automation [5]. Goldman Sachs projects long-term displacement of 6–7% of the U.S. workforce, where approximately 11 million workers and around 300 million full-time jobs are affected globally [6].

### **Positive Impacts of AI in Healthcare and Science**

AI has strong potential to improve healthcare by assisting doctors, reducing workload, and helping detect diseases earlier. In medical imaging, AI can analyze mammograms, X-rays, CT scans, and pathology images to identify patterns that may be difficult or time-consuming for humans to detect. For example, AI-assisted mammography has shown promise in

improving breast cancer detection while supporting, rather than replacing, medical professionals [16].

AI can also improve access to care in underserved communities. Screening tools for conditions such as diabetic retinopathy can help identify patients who need specialist attention, even in areas with limited medical resources. In addition, AI systems can help match patients to clinical trials more efficiently, allowing doctors to find treatment opportunities that may otherwise be missed [17].

Overall, AI's healthcare benefits are strongest when it functions as a decision-support tool. It can improve diagnosis, expand access, and reduce administrative burdens, but it must remain under human oversight to protect patient safety, privacy, and fairness.

AI is also accelerating scientific research by helping scientists analyze large datasets, generate hypotheses, and identify patterns more efficiently. One major example is AlphaFold, which predicts protein structures and has made millions of predictions available to researchers. This can support research in biology, medicine, and drug discovery by reducing the time and cost required for early scientific investigation [18].

AI can also assist in drug development by helping researchers identify possible drug candidates, predict molecular interactions, and prioritize experiments. Although laboratory testing and clinical trials are still necessary, AI can make the research process faster and more focused.

In this sense, AI does not replace scientists but expands their ability to explore complex problems. When used responsibly, AI can make scientific discovery more efficient, collaborative, and accessible.

## **Group Deliberation Process**

During our first group meeting on May 23 (length 13 minutes), the group reviewed the assignment requirements and identified the major areas of AI's societal impact that would be analyzed in the report. The group decided to focus on employment and labor markets, misuse through deepfakes and misinformation, privacy and surveillance, and positive impacts in

healthcare and science. Each member was assigned one area to research and was responsible for applying consequentialism, deontological ethics, and virtue ethics to that area.

During our final group meeting on May 25 (length 20 minutes), each member presented their findings and ethical analysis. The group compared the different sections and identified several common themes. One major point of agreement was that AI should not be evaluated only by its technical capabilities. Instead, the group found that ethical evaluation must consider who benefits, who is harmed, who is accountable, and whether the technology respects human rights and dignity.

A key tension emerged between consequentialist and deontological reasoning. In some areas, such as healthcare and scientific research, consequentialism highlighted AI's ability to produce major benefits, including earlier disease detection and faster discovery. However, deontology reminded the group that positive outcomes do not automatically justify systems that violate privacy, lack transparency, or treat individuals unfairly. Virtue ethics complemented both perspectives by emphasizing the responsibility, honesty, and humility required from developers, companies, governments, and users.

The group ultimately concluded that AI is not inherently ethical or unethical. Its moral value depends on the context of use, the quality of governance, and the values guiding its development. AI should be encouraged where it improves human well-being, but it must be limited or redesigned when it threatens fairness, privacy, accountability, or democratic trust.



## **Conclusion**

This report concludes that Artificial Intelligence has a complex ethical impact on society because it produces both significant benefits and serious risks. AI can improve productivity, support medical diagnosis, expand educational and scientific opportunities, and accelerate research. However, it can also displace workers, intensify inequality, enable misinformation, threaten privacy, and weaken accountability when used carelessly or without proper oversight.

The ethical frameworks used in this report show that no single perspective is sufficient by itself. Consequentialism is useful because it focuses on AI's real-world outcomes and helps compare benefits against harms. Deontological ethics is necessary because it protects rights, duties, and human dignity even when violating them might appear useful. Virtue ethics adds an important focus on the character and responsibility of the people and institutions developing AI.

Overall, the group's ethical evaluation is that AI should be developed and used under conditions of transparency, fairness, accountability, privacy protection, and meaningful human oversight. AI's positive applications, especially in healthcare and science, show that the technology can serve society when carefully governed. However, its risks show that technical progress alone is not enough. Ethical AI requires not only powerful systems, but responsible institutions and human-centered values.

## **Individual Reflections:**

### **Manuel:**

Working in the employment and labor section taught me that AI ethics is harder than picking a side. I learned that each framework works differently and catches different pros and cons between each perspective. Consequentialism showed me how unevenly AI's gains are disproportional to the amount of gains it can give in the labor force. Deontology made me aware of the issues caused by the algorithmic management case with Amazon, where it feels more ethically serious in a way that pure analysis could not really capture. Virtue ethics is the one I thought was more important: the impact in layoffs and indifference caused by viewing labor as nothing more than spreadsheets and results gained by AI while discarding human work made it show the negative impact AI can have on the workplace. I contributed by researching displacement data and helping apply the frameworks to the labor case, and knowing how one framework is not enough to deem AI's impact on society.

### **Emmanuel:**

In this project, I contributed by researching the positive impacts of AI in healthcare and science. I focused on how AI can improve medical diagnostics, support clinical-trial matching, expand access to screening, and accelerate scientific research through tools such as protein-structure prediction and AI-assisted data analysis. I learned that AI's benefits are strongest when the technology supports human experts rather than replacing them entirely. I also learned that ethical analysis is necessary even when AI has positive outcomes, because healthcare and science still involve privacy, fairness, safety, and accountability. My section helped balance the report by showing that AI is not only a source of risk, but also a technology with meaningful potential to improve human well-being when used responsibly.

### **Claudia:**

In this project, I focused on researching the ethical issues related to privacy, surveillance, and civil liberties in AI systems. I learned that AI technologies can create serious risks when they collect, analyze, or misuse personal data without transparency or consent. Through the ethical frameworks, I understood that consequentialism focuses on balancing security and social benefits against potential harm to individual freedoms, while deontological ethics emphasizes the importance of protecting privacy and respecting human rights regardless of outcomes. Virtue ethics helped me recognize the importance of responsibility, honesty, and accountability from governments and companies that develop or use AI systems. This project showed me that AI can be beneficial, but only if it is implemented with strong ethical safeguards and respect for human dignity.

**Carlos:**

In this project, I researched the ethical problems associated with AI-generated misinformation, deepfakes, and impersonation. While working on this section, I learned how rapidly AI can influence public opinion and spread false information in ways that are difficult to detect or control. Consequentialism helped me understand the large-scale social consequences that misinformation can produce, especially when it affects elections, public trust, or personal reputations. Deontological ethics showed me that creating or distributing deceptive AI content violates important duties related to honesty and respect for others. Virtue ethics also made me reflect on the importance of integrity and responsibility when developing or using AI technologies. Overall, this project helped me better understand that AI is not only a technical tool, but also a powerful social influence that requires ethical limits and responsible use.

## **Appendix:**

First meeting: May 23rd, 2026

Second meeting: May 25th, 2026

## **References**

[1] World Economic Forum, "The Future of Jobs Report 2023," Geneva, Switzerland, May 2023. [Online]. Available: <https://www.weforum.org/reports/the-future-of-jobs-report-2023>.

[2] K. Ellingrud *et al.*, "Generative AI and the future of work in America," McKinsey Global Institute, Jul. 2023. [Online]. Available: <https://www.mckinsey.com/mgi/our-research/generative-ai-and-the-future-of-work-in-america>.

[3] D. Shapiro, "AI destroyed 200k to 300k jobs in 2025 in the US," *Substack*, Feb. 2026. [Online]. Available: <https://daveshap.substack.com/p/ai-destroyed-200k-to-300k-jobs-in>

[4] *AI Job Displacement 2025: Which Jobs Are At Risk?*, FinalRound AI Research, 2025. [Online]. Available: <https://www.finalroundai.com/blog/ai-replacing-jobs-2025>

[5] World Economic Forum, *Future of Jobs Report 2025*, Geneva, Switzerland, 2025. [Online]. Available: <https://www.weforum.org/reports/the-future-of-jobs-report-2025>

[6] *AI Job Displacement Statistics (2025–2030): 60+ Facts on Which Jobs Are at Risk*, ALM Corp., Mar. 2026. [Online]. Available: <https://almcorp.com/blog/ai-job-displacement-statistics/>

[7] J. Hatzius *et al.*, "The potentially large effects of artificial intelligence on economic growth," Goldman Sachs Economics Research, Mar. 2023.

[8] D. Acemoglu, "The simple macroeconomics of AI," *Economic Policy*, vol. 40, no. 121, pp. 13–58, Jan. 2025.

[9] D. Acemoglu and S. Johnson, "Rebalancing AI," *Finance & Development*, International Monetary Fund, Dec. 2023. [Online]. Available: <https://www.imf.org/en/publications/fandd/issues/2023/12/rebalancing-ai-acemoglu-johnson>

- [10] J. Kim and S. Lee, "Artificial intelligence and wealth inequality: A comprehensive empirical exploration of socioeconomic implications," *Technology in Society*, 2024.
- [11] C. Lecher, "How Amazon automatically tracks and fires warehouse workers for 'productivity,'" *The Verge*, Apr. 2019. [Online]. Available: <https://www.theverge.com/2019/4/25/18516004/amazon-warehouse-fulfillment-centers-productivity-firing-terminations>
- [12] *Amazon AI Algorithm Fired 900 Employees: The Automation Crisis*, Yeet Magazine, 2026. [Online]. Available: <https://www.yeetmagazine.com/amazon-ai-algorithm-fired-900-employees/>
- [13] J. Kelly, "Everyone thinks AI is replacing factory workers, but Amazon's layoffs show it's coming for middle management first," *Fortune*, Oct. 29, 2025. [Online]. Available: <https://fortune.com/2025/10/29/amazon-layoffs-ai-middle-managers-robots-factory-workers/>
- [14] P. Wells, "Workday, Amazon AI employment bias claims add to growing concerns about the tech's hiring discrimination," *Fortune*, Jul. 5, 2025. [Online]. Available: <https://fortune.com/2025/07/05/workday-amazon-alleged-ai-employment-bias-hiring-discrimination/>
- [15] "AI & HR: Algorithmic discrimination in the workplace," *Cornell Journal of Law and Public Policy*, Nov. 2024. [Online]. Available: <https://publications.lawschool.cornell.edu/jlpp/2024/11/21/ai-hr-algorithmic-discrimination-in-the-workplace/>
- [16] N. Eisemann *et al.*, "Nationwide real-world implementation of AI for cancer detection in population-based mammography screening," *Nature Medicine*, vol. 31, pp. 917–924, Jan. 2025, doi: 10.1038/s41591-024-03408-6.
- [17] Q. Jin *et al.*, "Matching patients to clinical trials with large language models," *Nature Communications*, vol. 15, Art. no. 9074, Nov. 2024, doi: 10.1038/s41467-024-53081-z.